

## Week 4 - Homework

Alessio Benedetti

11 june 2018

### Question 9.2

Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2. You can use the R function `prcomp` for PCA.

First we need to load the libraries and the data from the temp `txt` file.

```
raw_data <- read.table('9.1uscrimeSummer2018.txt', stringsAsFactors = FALSE,
header=TRUE)
head(raw_data) #view top rows of dataset
```

##	M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq
## 1	15.1	1	9.1	5.8	5.6	0.510	95.0	33	30.1	0.108	4.1	3940	26.1
## 2	14.3	0	11.3	10.3	9.5	0.583	101.2	13	10.2	0.096	3.6	5570	19.4
## 3	14.2	1	8.9	4.5	4.4	0.533	96.9	18	21.9	0.094	3.3	3180	25.0
## 4	13.6	0	12.1	14.9	14.1	0.577	99.4	157	8.0	0.102	3.9	6730	16.7
## 5	14.1	0	12.1	10.9	10.1	0.591	98.5	18	3.0	0.091	2.0	5780	17.4
## 6	12.1	0	11.0	11.8	11.5	0.547	96.4	25	4.4	0.084	2.9	6890	12.6

##	Prob	Time	Crime
## 1	0.084602	26.2011	791
## 2	0.029599	25.2999	1635
## 3	0.083401	24.3006	578
## 4	0.015801	29.9012	1969
## 5	0.041399	21.2998	1234
## 6	0.034201	20.9995	682

As suggested in the headline, we can use the `prcomp` function to perform PCA with scaled values.

```
pca <- prcomp(raw_data[,1:15], scale=TRUE)
summary(pca)
```

## Importance of components:	PC1	PC2	PC3	PC4	PC5	PC6
## Standard deviation	2.4534	1.6739	1.4160	1.07806	0.97893	0.74377
## Proportion of Variance	0.4013	0.1868	0.1337	0.07748	0.06389	0.03688
## Cumulative Proportion	0.4013	0.5880	0.7217	0.79920	0.86308	0.89996


  

##	PC7	PC8	PC9	PC10	PC11	PC12
## Standard deviation	0.56729	0.55444	0.48493	0.44708	0.41915	0.35804

```
## Proportion of Variance 0.02145 0.02049 0.01568 0.01333 0.01171 0.00855
## Cumulative Proportion 0.92142 0.94191 0.95759 0.97091 0.98263 0.99117
##
##          PC13    PC14    PC15
## Standard deviation 0.26333 0.2418 0.06793
## Proportion of Variance 0.00462 0.0039 0.00031
## Cumulative Proportion 0.99579 0.9997 1.00000
```


The summary method returns the standard deviation of each of the principal components and their rotation. With the plot statement we're able to see the variances as a function of the principal components.

```
plot(pca, type = 'l')
```



We can see in the figure that the first 6 principal components explain around 90% of the variance of the data. We also can represent the biplot as shown below:

```
biplot(pca)
```



Now we can build the regression model by using the first 6 principal components. First we build a new dataframe with the 6 principal components and the response variable: Crime.

```
pca_df <- data.frame(cbind(pca$x[,1:6], raw_data$Crime))
names(pca_df) <- c('PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6', 'Crime')
```

Then we fit the regression model:

```
model_pca <- lm(Crime ~ ., pca_df)
summary(model_pca)
```

```
##
## Call:
## lm(formula = Crime ~ ., data = pca_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -377.15 -172.23   25.81  132.10  480.38
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   905.09      35.35   25.604 < 2e-16 ***
## PC1             65.22      14.56    4.478 6.14e-05 ***
## PC2            -70.08      21.35   -3.283 0.00214 **
## PC3             25.19      25.23    0.998 0.32409
## PC4             69.45      33.14    2.095 0.04252 *
## PC5            -229.04     36.50   -6.275 1.94e-07 ***
```

```
## PC6          -60.21      48.04  -1.253  0.21734
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 242.3 on 40 degrees of freedom
## Multiple R-squared:  0.6586, Adjusted R-squared:  0.6074
## F-statistic: 12.86 on 6 and 40 DF,  p-value: 4.869e-08
```

The R-squared value is about 65%. We should now convert the principal components of the model, back to the original factors by using the rotation matrix.

```
convert_coeff <- (pca$rotation[,1:6] %%%
model_pca$coefficients[2:7])/pca$scale
adjusted_intercept <- model_pca$coefficients[1] - sum(convert_coeff *
pca$center)
```

We the “rotated” model we can now predict the Crime response with the new model and compare it with the value obtained in exercise 8.2

```
new_datapoint <- data.frame(M = 14.0, So = 0, Ed = 10.0, Po1 = 12.0, Po2 =
15.5,
                           LF = 0.640, M.F = 94.0, Pop = 150, NW = 1.1, U1 = 0.120,
                           U2 = 3.6, Wealth = 3200, Ineq = 20.1, Prob = 0.04, Time
= 39.0)
```

```
Crime <- sum(
  convert_coeff[1,1] %%% new_datapoint$M,
  convert_coeff[2,1] %%% new_datapoint$So,
  convert_coeff[3,1] %%% new_datapoint$Ed,
  convert_coeff[4,1] %%% new_datapoint$Po1,
  convert_coeff[5,1] %%% new_datapoint$Po2,
  convert_coeff[6,1] %%% new_datapoint$LF,
  convert_coeff[7,1] %%% new_datapoint$M.F,
  convert_coeff[8,1] %%% new_datapoint$Pop,
  convert_coeff[9,1] %%% new_datapoint$NW,
  convert_coeff[10,1] %%% new_datapoint$U1,
  convert_coeff[11,1] %%% new_datapoint$U2,
  convert_coeff[12,1] %%% new_datapoint$Wealth,
  convert_coeff[13,1] %%% new_datapoint$Ineq,
  convert_coeff[14,1] %%% new_datapoint$Prob,
  convert_coeff[15,1] %%% new_datapoint$Time,
  adjusted_intercept
)
```

Crime

```
## [1] 1248.427
```

The Crime value for the new data point is 1248.427 with an R squared value of 65%.

In the previous exercise the regression model predicted a value of 1301.432 with an R squared value of 73%.

The PCA regression model performed a bit worse than the other.

### Question 10.1

**Using the same crime data set uscrime.txt as in Questions 8.2 and 9.1, find the best model you can using a regression tree model, and a random forest model. In R, you can use the tree package or the rpart package, and the randomForest package. For each model, describe one or two qualitative takeaways you get from analyzing the results**

The data is already loaded from the previous question, with only need to load the appropriate libraries to conduct our analysis.

```
#install.packages('tree')
library(tree)
#install.packages('randomForest')
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.
```


#### Tree

Let's first create a tree model.

```
model_tree <- tree(Crime ~ ., raw_data)
summary(model_tree)

##
## Regression tree:
## tree(formula = Crime ~ ., data = raw_data)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -98.300  -1.545    0.000 110.600  490.100

# Visualize tree model
plot(model_tree)
text(model_tree)
```



We can now calculate the R squared value (coefficient of determination) of the model. Let's build our own R squared function at first.

```
R2_calc <- function(yhat, raw_data) {
  SSres <- sum((yhat - raw_data$Crime)^2)
  SStot <- sum((raw_data$Crime - mean(raw_data$Crime))^2)
  R2 <- 1 - SSres/SStot
  return(R2)
}
```

Now we can calculate the coefficient of determination of the model:

```
crime_tree_yhat <- predict(model_tree)
crime_tree_r2 <- R2_calc(crime_tree_yhat, raw_data)
crime_tree_r2

## [1] 0.7244962
```

We calculated an R2 of near 72%.


We can prune the tree to prevent overfitting, by using the function `prune.tree()`. We only need to choose how many leafs we want the tree to have. Let's investigate with 3, 4 and 5 leaves.

```
model_tree_pruned_3 <- prune.tree(model_tree , best = 3)
summary(model_tree_pruned_3)
```

```
##
## Regression tree:
## snip.tree(tree = model_tree, nodes = c(6L, 2L, 7L))
## Variables actually used in tree construction:
## [1] "Po1" "NW"
## Number of terminal nodes: 3
## Residual mean deviance: 76460 = 3364000 / 44
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -550.9  -181.8   -37.9     0.0   158.9   688.1
```

```
plot(model_tree_pruned_3)
```

```
text(model_tree_pruned_3)
```



```
pruned_tree_yhat_3 <- predict(model_tree_pruned_3)
pruned_tree_3_r2 <- R2_calc(pruned_tree_yhat_3, raw_data)
pruned_tree_3_r2
```


```
## [1] 0.5111061
```

```
model_tree_pruned_4 <- prune.tree(model_tree , best = 4)
summary(model_tree_pruned_4)
```

```
##
## Regression tree:
## snip.tree(tree = model_tree, nodes = c(6L, 2L))
## Variables actually used in tree construction:
## [1] "Po1" "NW"
```

```
## Number of terminal nodes: 4
## Residual mean deviance: 61220 = 2633000 / 43
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.90 -152.60   35.39    0.00  158.90   490.10
```

```
plot(model_tree_pruned_4)
text(model_tree_pruned_4)
```



```
pruned_tree_yhat_4 <- predict(model_tree_pruned_4)
pruned_tree_4_r2 <- R2_calc(pruned_tree_yhat_4, raw_data)
pruned_tree_4_r2
```


```
## [1] 0.6174017
```

```
model_tree_pruned_5 <- prune.tree(model_tree , best = 5)
summary(model_tree_pruned_5)
```

```
##
## Regression tree:
## snip.tree(tree = model_tree, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  -573.9  -107.5   15.5     0.0   122.8   490.1
```



```
plot(model_tree_pruned_5)
text(model_tree_pruned_5)
```



```
pruned_tree_yhat_5 <- predict(model_tree_pruned_5)
pruned_tree_5_r2 <- R2_calc(pruned_tree_yhat_5, raw_data)
pruned_tree_5_r2
## [1] 0.6691333
```

As calculated the best model seems the tree pruned with 5 leaves: it gives an R squared of around 67 %.

## Random forest

Now we can use a random forest with 500 trees.

```
model_forest <- randomForest(Crime ~., raw_data, importance = TRUE, ntree =
500)
model_forest
##
## Call:
## randomForest(formula = Crime ~ ., data = raw_data, importance = TRUE,
ntree = 500)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 5
##
```

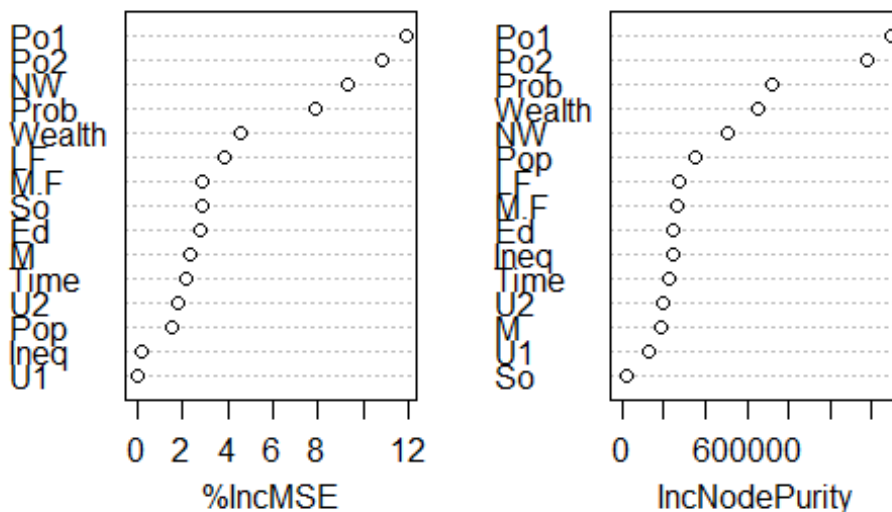
```
##           Mean of squared residuals: 82617.37
##           % Var explained: 43.57
```

```
importance(model_forest)
```

```
##           %IncMSE  IncNodePurity
## M           2.30706715    188669.79
## So          2.83705092     24824.93
## Ed          2.79500019    250519.88
## Po1         11.90074201   1292193.32
## Po2         10.84382663   1180597.40
## LF          3.90271783    274743.85
## M.F         2.86053480    263434.36
## Pop         1.55079227    353037.31
## NW          9.27667389    511844.38
## U1          -0.01202225    127143.07
## U2          1.84552751    196721.04
## Wealth      4.60664609    651559.62
## Ineq        0.23348414    244471.22
## Prob        7.88770708    725489.29
## Time        2.16628457    222730.25
```

```
varImpPlot(model_forest)
```

model\_forest



Once the model is created we can again evaluate it's R squared value

```
crime_forest_yhat <- predict(model_forest)
R2_calc(crime_forest_yhat, raw_data)
```

```
## [1] 0.4356842
```

The R squared value of the model is around 40%. The random forest has a lower R squared value, but this may come from the fact that the tree model (even pruned) has an overfitting component.

## Question 10.2

**Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.**

Since the football world cup is near, an example of logistic regression model may be the fact that a penalty succeeds (1) or fails (0). As predictors we can consider the level of fatigue of a player, power of kick, precision of kick, number of penalties he succeeded in his career and level of stress.

## Question 10.3.1

**Using the GermanCredit data set `germancredit.txt`, use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the `glm` function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use `family=binomial(link="logit")` in your `glm` function call.**

First we need to load the libraries and the data from the temp `txt` file.

```
raw_data <- read.table('10.3germancreditSummer2018.txt', stringsAsFactors = FALSE, header=FALSE)
head(raw_data) #view top rows of dataset
```

##	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14	V15	V16	V17
## 1	A11	6	A34	A43	1169	A65	A75	4	A93	A101	4	A121	67	A143	A152	2	A173
## 2	A12	48	A32	A43	5951	A61	A73	2	A92	A101	2	A121	22	A143	A152	1	A173
## 3	A14	12	A34	A46	2096	A61	A74	2	A93	A101	3	A121	49	A143	A152	1	A172
## 4	A11	42	A32	A42	7882	A61	A74	2	A93	A103	4	A122	45	A143	A153	1	A173
## 5	A11	24	A33	A40	4870	A61	A73	3	A93	A101	4	A124	53	A143	A153	2	A173
## 6	A14	36	A32	A46	9055	A65	A73	2	A93	A101	4	A124	35	A143	A153	1	A172
##	V18	V19	V20	V21													
## 1	1	A192	A201	1													
## 2	1	A191	A201	2													
## 3	2	A191	A201	1													
## 4	2	A191	A201	1													
## 5	2	A191	A201	2													
## 6	2	A192	A201	1													

We can initially change our response value in order to have zeroes (instead of 1) and ones (instead of 2)

```
raw_data$V21[raw_data$V21 == 1] <- 1
raw_data$V21[raw_data$V21 == 2] <- 0
```

Once done, we can fit a logistic model with all factors in order to determine the more significant (we split 80 / 20)

```
split <- sample(1:nrow(raw_data), size = round(0.8*(nrow(raw_data))))
train_df <- raw_data[split,]
test_df <- raw_data[-split,]

full_model <- glm(V21 ~ ., family = binomial(link="logit"), train_df)
summary(full_model)
```

```
##
## Call:
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = train_df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8198  -0.7129   0.3540   0.7101   2.2647
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.547e-01  1.205e+00  -0.294  0.768431
## V1A12        2.155e-01  2.471e-01   0.872  0.383135
## V1A13        3.242e-01  4.326e-01   0.749  0.453571
## V1A14        1.595e+00  2.637e-01   6.048  1.47e-09 ***
## V2          -3.224e-02  1.045e-02  -3.084  0.002039 **
## V3A31        2.275e-01  6.026e-01   0.378  0.705765
## V3A32        7.260e-01  4.757e-01   1.526  0.126960
## V3A33        1.300e+00  5.299e-01   2.453  0.014179 *
## V3A34        1.712e+00  4.939e-01   3.466  0.000528 ***
## V4A41        1.108e+00  4.144e-01   2.674  0.007495 **
## V4A410       1.880e+00  9.552e-01   1.968  0.049079 *
## V4A42        5.932e-01  2.939e-01   2.018  0.043554 *
## V4A43        7.756e-01  2.822e-01   2.749  0.005984 **
## V4A44        7.749e-01  8.608e-01   0.900  0.368017
## V4A45       -1.147e-01  6.246e-01  -0.184  0.854295
## V4A46       -5.749e-02  4.890e-01  -0.118  0.906408
## V4A48        1.710e+00  1.206e+00   1.418  0.156067
## V4A49        5.165e-01  3.753e-01   1.376  0.168821
## V5          -1.579e-04  5.004e-05  -3.156  0.001598 **
## V6A62        2.706e-01  3.097e-01   0.874  0.382388
## V6A63        7.410e-01  4.939e-01   1.500  0.133495
## V6A64        1.979e+00  6.744e-01   2.934  0.003346 **
## V6A65        1.002e+00  2.977e-01   3.366  0.000762 ***
## V7A72        4.700e-01  4.689e-01   1.002  0.316196
## V7A73        4.959e-01  4.537e-01   1.093  0.274324
## V7A74        1.352e+00  5.012e-01   2.696  0.007010 **
## V7A75        6.447e-01  4.577e-01   1.409  0.158895
```

```
## V8          -3.324e-01  1.021e-01  -3.256  0.001130 **
## V9A92       1.229e-01  4.407e-01   0.279  0.780341
## V9A93       6.406e-01  4.295e-01   1.492  0.135826
## V9A94       1.836e-01  5.133e-01   0.358  0.720603
## V10A102     -3.716e-01  4.291e-01  -0.866  0.386567
## V10A103     1.022e+00  4.673e-01   2.187  0.028768 *
## V11         -6.749e-02  9.899e-02  -0.682  0.495360
## V12A122     -3.294e-01  2.831e-01  -1.163  0.244735
## V12A123     -1.577e-01  2.630e-01  -0.600  0.548793
## V12A124     -6.843e-01  4.564e-01  -1.499  0.133806
## V13         1.449e-02  1.039e-02   1.395  0.163103
## V14A142     5.954e-01  4.673e-01   1.274  0.202565
## V14A143     7.995e-01  2.661e-01   3.005  0.002658 **
## V15A152     3.235e-01  2.694e-01   1.201  0.229809
## V15A153     9.699e-01  5.302e-01   1.829  0.067342 .
## V16         -2.349e-01  2.127e-01  -1.104  0.269394
## V17A172     -4.371e-01  7.294e-01  -0.599  0.549004
## V17A173     -5.854e-01  6.961e-01  -0.841  0.400402
## V17A174     -2.247e-01  7.037e-01  -0.319  0.749513
## V18         -2.707e-01  2.827e-01  -0.958  0.338292
## V19A192     3.493e-01  2.262e-01   1.544  0.122540
## V20A202     1.130e+00  6.601e-01   1.711  0.087043 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 970.51  on 799  degrees of freedom
## Residual deviance: 707.66  on 751  degrees of freedom
## AIC: 805.66
##
## Number of Fisher Scoring iterations: 5
```

We can now refine the model by selecting only the significant attributes, based on their p-values.

```
refined_model <- glm(V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V10 + V14 +
V20,
                    family = binomial(link="logit"), train_df
                    )
summary(refined_model)

##
## Call:
## glm(formula = V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V10 +
##      V14 + V20, family = binomial(link = "logit"), data = train_df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.6911  -0.7620   0.3929   0.7365   2.1340
```

```
##
## Coefficients:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) -4.702e-01  6.069e-01  -0.775 0.438444
## V1A12        2.150e-01  2.328e-01   0.923 0.355851
## V1A13        5.250e-01  4.025e-01   1.304 0.192087
## V1A14        1.612e+00  2.512e-01   6.417 1.39e-10 ***
## V2          -3.023e-02  9.743e-03  -3.103 0.001915 **
## V3A31        5.441e-01  5.736e-01   0.949 0.342828
## V3A32        9.476e-01  4.454e-01   2.127 0.033401 *
## V3A33        1.457e+00  5.132e-01   2.839 0.004526 **
## V3A34        1.837e+00  4.710e-01   3.900 9.62e-05 ***
## V4A41        1.070e+00  3.922e-01   2.728 0.006381 **
## V4A410       1.759e+00  8.660e-01   2.031 0.042278 *
## V4A42        3.857e-01  2.748e-01   1.404 0.160431
## V4A43        7.255e-01  2.676e-01   2.711 0.006706 **
## V4A44        6.400e-01  8.023e-01   0.798 0.425041
## V4A45       -2.228e-01  6.117e-01  -0.364 0.715701
## V4A46       -1.280e-01  4.662e-01  -0.275 0.783624
## V4A48        1.744e+00  1.188e+00   1.467 0.142289
## V4A49        6.071e-01  3.571e-01   1.700 0.089119 .
## V5          -1.048e-04  4.441e-05  -2.360 0.018282 *
## V6A62        1.808e-01  2.920e-01   0.619 0.535647
## V6A63        8.661e-01  4.769e-01   1.816 0.069338 .
## V6A64        1.889e+00  6.617e-01   2.855 0.004301 **
## V6A65        9.844e-01  2.805e-01   3.509 0.000449 ***
## V8          -2.483e-01  9.486e-02  -2.617 0.008867 **
## V10A102      -4.342e-01  4.180e-01  -1.039 0.298847
## V10A103      9.917e-01  4.379e-01   2.265 0.023524 *
## V14A142      4.907e-01  4.484e-01   1.094 0.273845
## V14A143      7.118e-01  2.548e-01   2.793 0.005224 **
## V20A202      1.026e+00  6.299e-01   1.629 0.103319
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 970.51  on 799  degrees of freedom
## Residual deviance: 744.46  on 771  degrees of freedom
## AIC: 802.46
##
## Number of Fisher Scoring iterations: 5
```

Now we can evaluate the refined model by calculating the confusion matrix. We added a step for converting the continuous output in a 0/1 response.

```
refined_model_yhat <- predict(refined_model, test_df, type = "response")
norm_refined_model_yhat <- as.integer(refined_model_yhat > 0.5)
table(test_df$V21, norm_refined_model_yhat)
```

```
##      norm_refined_model_yhat
##      0      1
##    0  29  35
##    1  16 120
```

Based on the resulting confusion matrix (row = true classification & col = model's classification) we got 29 false negatives and 20 false positives.

```
sensitivity = 28/(28+29)
specificity = 123/(20+123)
```

With a sensitivity = 0.4912281 and a specificity = 0.8601399

### Question 10.3.2

**Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between “good” and “bad” answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.**

In other terms false positives (incorrectly identifying a bad customer as good) are 5 times worse than false negatives (incorrectly classifying a good customer as bad).

We may work on the 50% level used before to convert the continuous output in a 0/1 response. By looping from 8% to 99% we can calculate each time a cost function as  $FN + 5*FP$ .

```
for (level in 8:99) {
  norm_refined_model_yhat <- as.integer(refined_model_yhat > level/100)
  table_level <- table(test_df$V21, norm_refined_model_yhat)
  cost_level <- table_level[1,2] + 5*table_level[2,1]
  print (paste(level,cost_level))
}

## [1] "8 63"
## [1] "9 63"
## [1] "10 62"
## [1] "11 62"
## [1] "12 61"
## [1] "13 60"
## [1] "14 60"
## [1] "15 60"
## [1] "16 59"
## [1] "17 59"
## [1] "18 59"
## [1] "19 59"
## [1] "20 59"
## [1] "21 58"
## [1] "22 58"
## [1] "23 58"
```

```
## [1] "24 57"  
## [1] "25 60"  
## [1] "26 59"  
## [1] "27 64"  
## [1] "28 62"  
## [1] "29 61"  
## [1] "30 59"  
## [1] "31 58"  
## [1] "32 57"  
## [1] "33 60"  
## [1] "34 59"  
## [1] "35 59"  
## [1] "36 63"  
## [1] "37 63"  
## [1] "38 68"  
## [1] "39 68"  
## [1] "40 68"  
## [1] "41 72"  
## [1] "42 72"  
## [1] "43 76"  
## [1] "44 75"  
## [1] "45 74"  
## [1] "46 79"  
## [1] "47 84"  
## [1] "48 99"  
## [1] "49 117"  
## [1] "50 115"  
## [1] "51 119"  
## [1] "52 119"  
## [1] "53 138"  
## [1] "54 137"  
## [1] "55 135"  
## [1] "56 135"  
## [1] "57 135"  
## [1] "58 144"  
## [1] "59 148"  
## [1] "60 152"  
## [1] "61 161"  
## [1] "62 166"  
## [1] "63 164"  
## [1] "64 164"  
## [1] "65 163"  
## [1] "66 173"  
## [1] "67 175"  
## [1] "68 180"  
## [1] "69 184"  
## [1] "70 183"  
## [1] "71 183"  
## [1] "72 192"  
## [1] "73 202"
```



```
## [1] "74 212"  
## [1] "75 227"  
## [1] "76 247"  
## [1] "77 262"  
## [1] "78 280"  
## [1] "79 285"  
## [1] "80 294"  
## [1] "81 307"  
## [1] "82 316"  
## [1] "83 316"  
## [1] "84 329"  
## [1] "85 353"  
## [1] "86 387"  
## [1] "87 402"  
## [1] "88 401"  
## [1] "89 420"  
## [1] "90 435"  
## [1] "91 455"  
## [1] "92 490"  
## [1] "93 528"  
## [1] "94 557"  
## [1] "95 572"  
## [1] "96 607"  
## [1] "97 625"  
## [1] "98 635"  
## [1] "99 670"
```

By observing the minimum (53) and maximum values (695) of the cost value we can suggest a threshold at 29% to be under the 10% of the maximum cost value.