

Homework 4 - Week 4

Akylas Stratigakos

12 Jun 2018

Question 9.1

In this task a linear regression model will be applied on the “uscrime” data, after reducing the dimensions of the predictors via Principal Component Analysis. Besides dimensionality reduction, PCA will also ensure no collinearity between the predictors. First, I will apply PCA on the dataset (without the response column, i.e. crime). In order to select a reasonable number of PCs, I plot the cumulative variance explained by the PCs (which depends on the eigenvalues of each component), and will select a number that explains around 80% of the total variance in the predictors.

```
datapath=paste("C:/Users/a.stratigakos/Desktop/edx/Introduction to Analytics Modelling"
               ,"/Week 2/UScrime_data.txt",sep="")
uscrime<-(read.table(datapath,header =TRUE))

#PCA without the response variable
PC=prcomp(uscrime[,-16], center = TRUE, scale. = TRUE)
summary(PC)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	2.4534	1.6739	1.4160	1.07806	0.97893	0.74377
Proportion of Variance	0.4013	0.1868	0.1337	0.07748	0.06389	0.03688
Cumulative Proportion	0.4013	0.5880	0.7217	0.79920	0.86308	0.89996

	PC7	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.56729	0.55444	0.48493	0.44708	0.41915	0.35804
Proportion of Variance	0.02145	0.02049	0.01568	0.01333	0.01171	0.00855
Cumulative Proportion	0.92142	0.94191	0.95759	0.97091	0.98263	0.99117


	PC13	PC14	PC15
Standard deviation	0.26333	0.2418	0.06793
Proportion of Variance	0.00462	0.0039	0.00031
Cumulative Proportion	0.99579	0.9997	1.00000

```
Threshold=matrix(0.8,nrow=16)
VarExpl=cumsum((PC$sdev)^2)/sum(PC$sdev^2)
plot(VarExpl,main='Cumulative Proportion of variance explained',
     ylab='Principal Components',
     xlab="Percentage")
lines(Threshold,type="l",col="red")
```

Cumulative Proportion of variance explained



```
biplot(PC,choices = 1:2,scale = 1)
```



```
PCcoeff=PC$rotation
scores=PC$x
```

As can be seen above, the first 4 Components explain around 80% of the variance, so I'll select them as the predictors. Additionally, the biplot of the first 2 PCs is shown, which shows the **coefficients** (eigenvectors or weights) of each variable and it's usefull for qualitatevely analysis. Afterwards, I proceed with applying a linear regression using the Crime collumn as response and the **scores** (i.e. original data points projected on the new dimensions) as the predictors. This methodology will be referred to as "PCA regression".

```
Crime=uscrime[,16];Sc1=scores[,1];Sc2=scores[,2]
Sc3=scores[,3];Sc4=scores[,4]
PCdf=data.frame(Crime,Sc1,Sc2,Sc3,Sc4)
PCmodel=lm(Crime~.,data=PCdf)
summary(PCmodel)
```

Call:

```
lm(formula = Crime ~ ., data = PCdf)
```

Residuals:

Min	1Q	Median	3Q	Max
-557.76	-210.91	-29.08	197.26	810.35

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	905.09	49.07	18.443	< 2e-16 ***
Sc1	65.22	20.22	3.225	0.00244 **

```
Sc2      -70.08      29.63    -2.365    0.02273 *
Sc3       25.19      35.03     0.719    0.47602
Sc4       69.45      46.01     1.509    0.13872
```

```
---
```

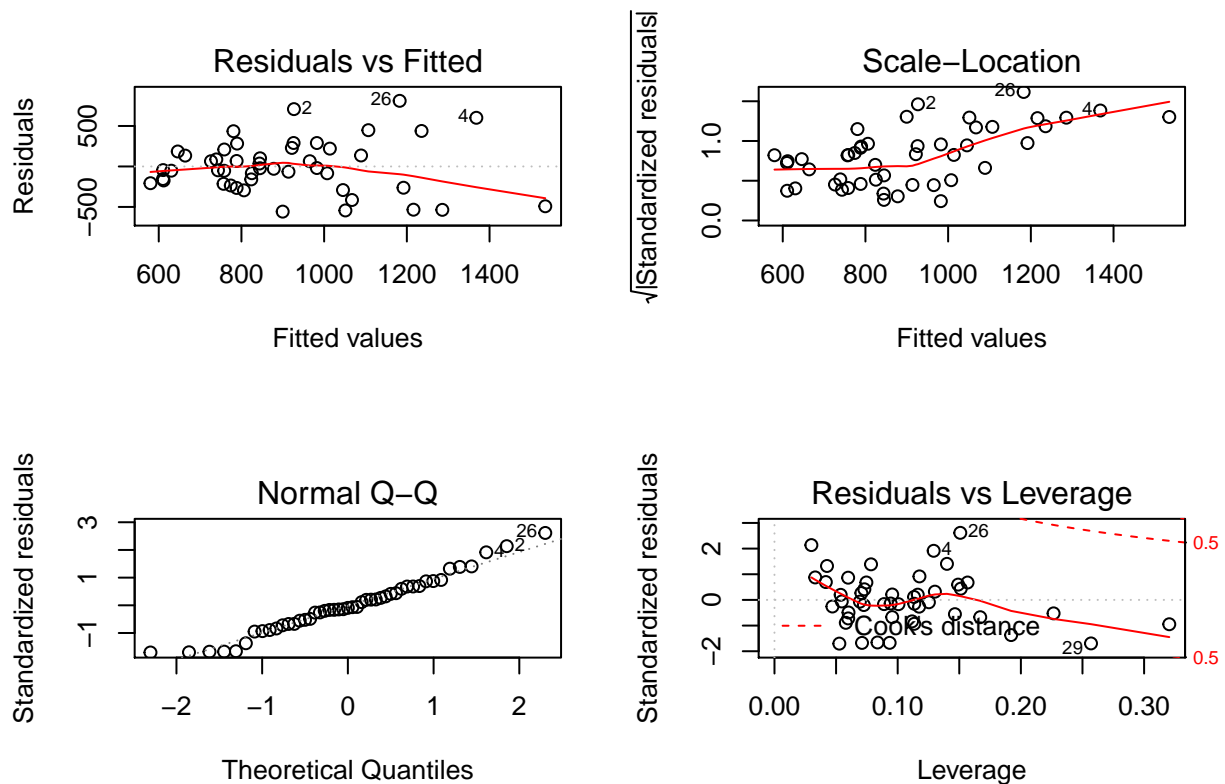
```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 336.4 on 42 degrees of freedom
```

```
Multiple R-squared:  0.3091,    Adjusted R-squared:  0.2433
```

```
F-statistic: 4.698 on 4 and 42 DF,  p-value: 0.003178
```

```
layout(matrix(c(1,2,3,4),2,2)) # optional 4 graphs/page
plot(PCmodel)
```



```
a=summary(PCmodel)$coefficients
Rsqr=summary(PCmodel)$r.squared
adjRsqr=summary(PCmodel)$adj.r.squared
Fstat=summary(PCmodel)$fstatistic
```

Above we can see some diagnostics about the PCA regression such as F-statistic, R squared and adjusted R square. Since PCA is a linear combination of the original dataset, the PCA regression will be also be this transformation multiplied by the regression's coefficients, which in turn results again in a linear combination of the original variables. Note the variables **scaled**. Due to space limitations I will only show the first coefficient and first Principal Component in the equation:

$$Y = 905.0851064 + 65.2159301 * X_{np} * \begin{pmatrix} -0.3037119 \\ -0.3308813 \\ 0.3396215 \\ 0.3086341 \\ 0.3109929 \\ 0.1761776 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \\ -0.3037119 \end{pmatrix} + \dots$$

where:

$$X = (M \quad So \quad Ed \quad Po1 \quad \dots \quad Ineq \quad Prob \quad Time)$$

I also fitted a simple linear regression model, with the whole dataset as a predictor and plotted the diagnostics. Afterwards I used the two models to make a prediction based on the new datapoint given in Question 8.2. Note that in order to make a prediction with the PCA regression I need to map the new data point into the Principal Components. After **scaling** them, I use matrix multiplication with the corresponding eigenvectors (coefficients), and then use them for prediction.

```
#Simple linear reg
fit <- lm( Crime~., data=uscrime)
summary(fit)
```

Call:

```
lm(formula = Crime ~ ., data = uscrime)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-395.74	-98.09	-6.69	112.99	512.67

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-5.984e+03	1.628e+03	-3.675	0.000893	***
M	8.783e+01	4.171e+01	2.106	0.043443	*
So	-3.803e+00	1.488e+02	-0.026	0.979765	
Ed	1.883e+02	6.209e+01	3.033	0.004861	**
Po1	1.928e+02	1.061e+02	1.817	0.078892	.
Po2	-1.094e+02	1.175e+02	-0.931	0.358830	
LF	-6.638e+02	1.470e+03	-0.452	0.654654	
M.F	1.741e+01	2.035e+01	0.855	0.398995	
Pop	-7.330e-01	1.290e+00	-0.568	0.573845	
NW	4.204e+00	6.481e+00	0.649	0.521279	
U1	-5.827e+03	4.210e+03	-1.384	0.176238	
U2	1.678e+02	8.234e+01	2.038	0.050161	.

```

Wealth      9.617e-02  1.037e-01  0.928 0.360754
Ineq        7.067e+01  2.272e+01  3.111 0.003983 **
Prob       -4.855e+03  2.272e+03  -2.137 0.040627 *
Time       -3.479e+00  7.165e+00  -0.486 0.630708
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

Residual standard error: 209.1 on 31 degrees of freedom
Multiple R-squared:  0.8031,    Adjusted R-squared:  0.7078
F-statistic: 8.429 on 15 and 31 DF,  p-value: 3.539e-07

```

```

lmRsqr=summary(fit)$r.squared
lmadjRsqr=summary(fit)$adj.r.squared
lmFstat=summary(fit)$fstatistic;
#Predict
M = 14.0;So = 0;Ed = 10.0;Po1 = 12.0;Po2 = 15.5;LF = 0.640
M.F = 94.0;Pop = 150;NW = 1.1;U1 = 0.120;U2 = 3.6;Wealth = 3200
Ineq = 20.1;Prob = 0.04;Time = 39.0
new.df <- data.frame(1,M,So, Ed, Po1, Po2, LF, M.F, Pop, NW, U1, U2, Wealth, Ineq, Prob, Time)
y=matrix(c(M,So, Ed, Po1, Po2, LF, M.F, Pop, NW, U1, U2, Wealth, Ineq, Prob, Time),nrow=1,ncol=15)
#Map data frame into the first 2 PC
newScores=((y-PC$center)/t(PC$scale))%*%PC$coeff[,1:4]
new.PCdf <- data.frame(1,Sc1=newScores[,1],Sc2=newScores[,2],Sc3=newScores[,3],Sc4=newScores[,4])
PCpred=predict(PCmodel, new.PCdf) #Scaling needs fixing
#LM pred
LMpred=predict(fit, new.df)
PCpred

```

```

      PC1
1112.678
LMpred

```

```

      1
155.4349

```

Question 10.1

In this Question the focus is to find a good model for the “uscrime” data using (a) a regression tree and (b) a random forest, while also provide some qualitative takeaways from analyzing the results.

Part a: First, a regression tree model will be fitted in the dataset, by using the standard recursive binary split. For this purpose the **tree** package is utilized.

```

datapath=paste("C:/Users/a.stratigakos/Desktop/edx/Introduction to Analytics Modelling"
               ,"/Week 2/UScrime_data.txt",sep="")
uscrime<-(read.table(datapath,header =TRUE))
library(rpart)
rtree=rpart(Crime~., data=uscrime, method="anova")
printcp(rtree) # display the results

```

```

Regression tree:
rpart(formula = Crime ~ ., data = uscrime, method = "anova")

```

```

Variables actually used in tree construction:
[1] NW  Po1 Pop


```

Root node error: 6880928/47 = 146403

n= 47

	CP	nsplit	rel error	xerror	xstd
1	0.362963	0	1.00000	1.06854	0.26591
2	0.148143	1	0.63704	1.05518	0.22490
3	0.051732	2	0.48889	1.00648	0.21263
4	0.010000	3	0.43716	0.93739	0.19926

```
plotcp(rtree) # visualize cross-validation results
```



```
summary(rtree) # detailed summary of splits
```

Call:

```
rpart(formula = Crime ~ ., data = uscrime, method = "anova")  
n= 47
```

	CP	nsplit	rel error	xerror	xstd
1	0.36296293	0	1.0000000	1.0685435	0.2659092
2	0.14814320	1	0.6370371	1.0551764	0.2248958
3	0.05173165	2	0.4888939	1.0064809	0.2126293
4	0.01000000	3	0.4371622	0.9373906	0.1992577

Variable importance

Po1	Po2	Wealth	Ineq	Prob	M	NW	Pop	Time	Ed
-----	-----	--------	------	------	---	----	-----	------	----

17	17	11	11	10	10	9	5	4	4
LF	So								
1	1								

Node number 1: 47 observations, complexity param=0.3629629

mean=905.0851, MSE=146402.7

left son=2 (23 obs) right son=3 (24 obs)

Primary splits:

Po1	< 7.65	to the left,	improve=0.3629629, (0 missing)
Po2	< 7.2	to the left,	improve=0.3629629, (0 missing)
Prob	< 0.0418485	to the right,	improve=0.3217700, (0 missing)
NW	< 7.65	to the left,	improve=0.2356621, (0 missing)
Wealth	< 6240	to the left,	improve=0.2002403, (0 missing)

Surrogate splits:

Po2	< 7.2	to the left,	agree=1.000, adj=1.000, (0 split)
Wealth	< 5330	to the left,	agree=0.830, adj=0.652, (0 split)
Prob	< 0.043598	to the right,	agree=0.809, adj=0.609, (0 split)
M	< 13.25	to the right,	agree=0.745, adj=0.478, (0 split)
Ineq	< 17.15	to the right,	agree=0.745, adj=0.478, (0 split)

Node number 2: 23 observations, complexity param=0.05173165

mean=669.6087, MSE=33880.15

left son=4 (12 obs) right son=5 (11 obs)

Primary splits:

Pop	< 22.5	to the left,	improve=0.4568043, (0 missing)
M	< 14.5	to the left,	improve=0.3931567, (0 missing)
NW	< 5.4	to the left,	improve=0.3184074, (0 missing)
Po1	< 5.75	to the left,	improve=0.2310098, (0 missing)
U1	< 0.093	to the right,	improve=0.2119062, (0 missing)

Surrogate splits:

NW	< 5.4	to the left,	agree=0.826, adj=0.636, (0 split)
M	< 14.5	to the left,	agree=0.783, adj=0.545, (0 split)
Time	< 22.30055	to the left,	agree=0.783, adj=0.545, (0 split)
So	< 0.5	to the left,	agree=0.739, adj=0.455, (0 split)
Ed	< 10.85	to the right,	agree=0.739, adj=0.455, (0 split)

Node number 3: 24 observations, complexity param=0.1481432

mean=1130.75, MSE=150173.4

left son=6 (10 obs) right son=7 (14 obs)

Primary splits:

NW	< 7.65	to the left,	improve=0.2828293, (0 missing)
M	< 13.05	to the left,	improve=0.2714159, (0 missing)
Time	< 21.9001	to the left,	improve=0.2060170, (0 missing)
M.F	< 99.2	to the left,	improve=0.1703438, (0 missing)
Po1	< 10.75	to the left,	improve=0.1659433, (0 missing)

Surrogate splits:

Ed	< 11.45	to the right,	agree=0.750, adj=0.4, (0 split)
Ineq	< 16.25	to the left,	agree=0.750, adj=0.4, (0 split)
Time	< 21.9001	to the left,	agree=0.750, adj=0.4, (0 split)
Pop	< 30	to the left,	agree=0.708, adj=0.3, (0 split)
LF	< 0.5885	to the right,	agree=0.667, adj=0.2, (0 split)

Node number 4: 12 observations

mean=550.5, MSE=20317.58

Node number 5: 11 observations
mean=799.5455, MSE=16315.52

Node number 6: 10 observations
mean=886.9, MSE=55757.49

Node number 7: 14 observations
mean=1304.929, MSE=144801.8

```
# create additional plots
par(mfrow=c(1,2)) # two plots on one page
rsq.rpart(rtree) # visualize cross-validation results
```

Regression tree:

```
rpart(formula = Crime ~ ., data = uscrime, method = "anova")
```

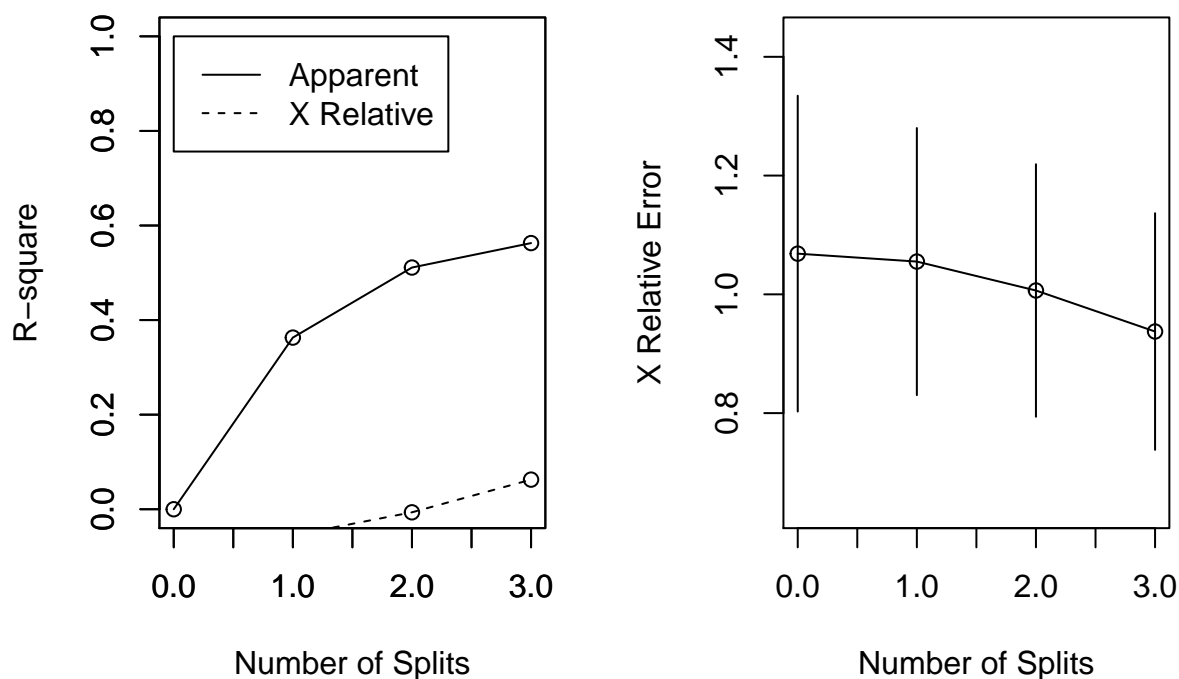
Variables actually used in tree construction:

```
[1] NW  Po1 Pop
```

Root node error: 6880928/47 = 146403

n= 47


	CP	nsplit	rel error	xerror	xstd
1	0.362963	0	1.00000	1.06854	0.26591
2	0.148143	1	0.63704	1.05518	0.22490
3	0.051732	2	0.48889	1.00648	0.21263
4	0.010000	3	0.43716	0.93739	0.19926




The function automatically selects the regression tree which minimizes the cross-validation error. A variety of diagnostics such as the cross-validation results vs number splits, importance of the variables and the improvement of the model for each split are plotted above. The most important are Police expenditure in 1960, expenditure in 1959, wealth and income inequality. We can plot the regression tree in order to visualize the results better. Additionally, we can prune the tree in order to avoid overfitting the data.

```
pfit<- prune(rtree, cp=0.01160389) # from cptable
# plot the pruned tree
par(mfrow=c(1,2)) # two plots on one page
plot(rtree, uniform=TRUE,
     main="Regression Tree for Crime")
text(rtree, use.n=TRUE, all=TRUE, cex=.8)
plot(pfit, uniform=TRUE,
     main="Pruned Regression Tree for Crime")
text(rtree, use.n=TRUE, all=TRUE, cex=.8)
```

Regression Tree for Crime



Pruned Regression Tree for Crim



The prune method results in the same model as the initial regression. The first split is made for **Police expenditure in 1960 < 7.65**. If this inequality is true, the second split is made for **State Population < 22.5**, while if its not true we check whether the **Number of non-whites per 1000 people** is < 7.65. Overall the tree has 3 nodes, and 4 leaves.

Part b: Now we apply a random forest algorithm in the dataset. Generally, random forests improve significantly over the simple regression trees, but they are harder to interpret. The **randomForest** package was used.

```
library(randomForest)
rf<- randomForest(Crime~.,data=uscrime,ntree=1000)
print(rf) # view results
```

Call:

```
randomForest(formula = Crime ~ ., data = uscrime, ntree = 1000)
      Type of random forest: regression
      Number of trees: 1000
```

No. of variables tried at each split: 5


```
      Mean of squared residuals: 82462.27
      % Var explained: 43.67
```

```
importance(rf) # importance of each predictor
```

```
      IncNodePurity
M      204293.55
So     18358.69
```

Ed	258121.29
Po1	1234820.69
Po2	1072613.32
LF	270494.90
M.F	274824.77
Pop	335363.36
NW	536890.06
U1	138162.18
U2	175420.12
Wealth	692032.60
Ineq	213707.90
Prob	712482.08
Time	222096.78

```
plot(rf,main="Random Forest")
```



In the plot above we can see the relative error vs the number of trees used. For each tree 5 predictors were utilized. The **importance** function shows that the most important variables are Police expenditure in 1960, Police expenditure in 1959, probability of imprisonment and wealth.

Question 10.2

Logistic regression is a go-to method for binary classification problems and for calculating probability of events. For example it could be used to calculate the probability of penalty kick success in soccer. List of predictors could include height of the player, whether he uses right or left foot to shoot the ball, strength of his kick, arm span of the goalkeeper etc.

Question 10.3

Part 1: In this task I'll use a logistic regression in order to find a good predictive model for whether credit applicants are “good” or “bad”. I define as positive state in my model the “**bad**” state. First I'll load the data and do some manipulation in order to produce more presentable results (create binary output by setting “good” as 0 and “bad” as 1). I use a 80/20 split in my dataset, for training and test set. I use the package **caret** in order to produce the confusion matrix and some helpfull diagnostics, such as accuracy, sensitivity and specificity. I'll also plot some diagnostics of the logistic regression such as z-stat of the coefficients.

```
datapath=paste("C:/Users/a.stratigakos/Desktop/edx/Introduction to Analytics Modelling"
               ,"/Week 4/German_credit.txt",sep="")
library(caret)
German_credit<-(read.table(datapath,header =TRUE))
str(German_credit)
```

```
'data.frame':  999 obs. of  21 variables:
 $ A11  : Factor w/ 4 levels "A11","A12","A13",...: 2 4 1 1 4 4 2 4 2 2 ...
 $ X6   : int   48 12 42 24 36 24 36 12 30 12 ...
 $ A34  : Factor w/ 5 levels "A30","A31","A32",...: 3 5 3 4 3 3 3 3 5 3 ...
 $ A43  : Factor w/ 10 levels "A40","A41","A410",...: 5 8 4 1 8 4 2 5 1 1 ...
 $ X1169: int  5951 2096 7882 4870 9055 2835 6948 3059 5234 1295 ...
 $ A65  : Factor w/ 5 levels "A61","A62","A63",...: 1 1 1 1 5 3 1 4 1 1 ...
 $ A75  : Factor w/ 5 levels "A71","A72","A73",...: 3 4 4 3 3 5 3 4 1 2 ...
 $ X4   : int   2 2 2 3 2 3 2 2 4 3 ...
 $ A93  : Factor w/ 4 levels "A91","A92","A93",...: 2 3 3 3 3 3 3 1 4 2 ...
 $ A101 : Factor w/ 3 levels "A101","A102",...: 1 1 3 1 1 1 1 1 1 1 ...
 $ X4.1 : int   2 3 4 4 4 4 2 4 2 1 ...
 $ A121 : Factor w/ 4 levels "A121","A122",...: 1 1 2 4 4 2 3 1 3 3 ...
 $ X67  : int   22 49 45 53 35 53 35 61 28 25 ...
 $ A143 : Factor w/ 3 levels "A141","A142",...: 3 3 3 3 3 3 3 3 3 3 ...
 $ A152 : Factor w/ 3 levels "A151","A152",...: 2 2 3 3 3 2 1 2 2 1 ...
 $ X2   : int   1 1 1 2 1 1 1 1 2 1 ...
 $ A173 : Factor w/ 4 levels "A171","A172",...: 3 2 3 3 2 3 4 2 4 3 ...
 $ X1   : int   1 2 2 2 2 1 1 1 1 1 ...
 $ A192 : Factor w/ 2 levels "A191","A192": 1 1 1 1 2 1 2 1 1 1 ...
 $ A201 : Factor w/ 2 levels "A201","A202": 1 1 1 1 1 1 1 1 1 1 ...
 $ X1.1 : int   2 1 1 2 1 1 1 1 2 2 ...
```

```
Actual<-ifelse(German_credit[,21]==1,"GOOD", "BAD")
German_credit[which(German_credit[,21]==1),21]=0 #GOOD
German_credit[which(German_credit[,21]==2),21]=1 #BAD
spec = c(train = .80, test = .20)
g = sample(cut(seq(nrow(German_credit)), nrow(German_credit)*cumsum(c(0,spec))),labels = names(spec)))
traind=as.data.frame(German_credit[g=="train",])
testd=as.data.frame(German_credit[g=="test",])
#transform to X=0 GOOD, X=1 BAD
#Predictive model with logistic regression
logitMod <- glm(X1.1~., data=traind, family=binomial(link="logit"))
summary(logitMod)
```

Call:

```
glm(formula = X1.1 ~ ., family = binomial(link = "logit"), data = traind)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.2575	-0.6780	-0.3515	0.6963	2.6243

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	7.630e-01	1.240e+00	0.615	0.538278	
A11A12	-5.980e-01	2.492e-01	-2.399	0.016420	*
A11A13	-1.268e+00	4.339e-01	-2.923	0.003471	**
A11A14	-1.891e+00	2.667e-01	-7.092	1.32e-12	***
X6	2.505e-02	1.053e-02	2.379	0.017337	*
A34A31	2.680e-01	6.256e-01	0.428	0.668424	
A34A32	-5.607e-01	5.029e-01	-1.115	0.264931	
A34A33	-4.490e-01	5.396e-01	-0.832	0.405387	
A34A34	-1.421e+00	5.091e-01	-2.791	0.005261	**
A43A41	-1.614e+00	4.162e-01	-3.878	0.000105	***
A43A410	-1.901e+00	8.391e-01	-2.265	0.023497	*
A43A42	-9.697e-01	2.948e-01	-3.289	0.001004	**
A43A43	-1.018e+00	2.853e-01	-3.568	0.000359	***
A43A44	-6.799e-01	8.065e-01	-0.843	0.399214	
A43A45	1.809e-02	6.082e-01	0.030	0.976273	
A43A46	-2.735e-01	4.528e-01	-0.604	0.545860	
A43A48	-2.080e+00	1.306e+00	-1.592	0.111307	
A43A49	-7.535e-01	3.797e-01	-1.984	0.047220	*
X1169	1.213e-04	4.969e-05	2.441	0.014627	*
A65A62	-2.976e-01	3.180e-01	-0.936	0.349359	
A65A63	-4.643e-01	4.299e-01	-1.080	0.280094	
A65A64	-1.456e+00	5.937e-01	-2.453	0.014173	*
A65A65	-1.097e+00	3.108e-01	-3.529	0.000417	***
A75A72	2.139e-01	4.715e-01	0.454	0.650027	
A75A73	3.307e-01	4.481e-01	0.738	0.460518	
A75A74	-5.892e-01	4.838e-01	-1.218	0.223213	
A75A75	-8.321e-02	4.533e-01	-0.184	0.854350	
X4	3.515e-01	1.016e-01	3.460	0.000540	***
A93A92	-7.354e-02	4.294e-01	-0.171	0.864030	
A93A93	-4.435e-01	4.191e-01	-1.058	0.289925	
A93A94	-2.787e-01	5.258e-01	-0.530	0.596049	
A101A102	3.098e-01	4.678e-01	0.662	0.507772	
A101A103	-1.097e+00	5.004e-01	-2.191	0.028437	*
X4.1	5.084e-02	9.911e-02	0.513	0.607978	
A121A122	2.133e-01	2.928e-01	0.728	0.466371	
A121A123	2.344e-01	2.702e-01	0.867	0.385786	
A121A124	8.491e-01	4.776e-01	1.778	0.075432	.
X67	-1.367e-02	1.049e-02	-1.303	0.192427	
A143A142	-6.222e-01	4.828e-01	-1.289	0.197547	
A143A143	-8.858e-01	2.736e-01	-3.238	0.001204	**
A152A152	-4.019e-01	2.681e-01	-1.499	0.133877	
A152A153	-6.747e-01	5.389e-01	-1.252	0.210574	
X2	1.571e-01	2.256e-01	0.697	0.486016	
A173A172	3.819e-01	7.684e-01	0.497	0.619145	
A173A173	3.090e-01	7.359e-01	0.420	0.674548	
A173A174	4.466e-01	7.625e-01	0.586	0.558038	
X1	-2.826e-02	2.919e-01	-0.097	0.922868	
A192A192	-3.959e-01	2.305e-01	-1.717	0.085902	.
A201A202	-1.224e+00	6.569e-01	-1.864	0.062358	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 980.03 on 798 degrees of freedom
Residual deviance: 700.80 on 750 degrees of freedom
AIC: 798.8

Number of Fisher Scoring iterations: 5

```
predicted <- predict(logitMod, testd, type="response") # predicted scores
Prediction=ifelse(predicted>0.5,'BAD','GOOD')
cM=confusionMatrix(table(Prediction,Actual[g=="test"]))
confusionMatrix(table(Prediction,Actual[g=="test"]))
```

Confusion Matrix and Statistics

Prediction BAD GOOD

BAD	26	22
GOOD	32	120

Accuracy : 0.73
95% CI : (0.6628, 0.7902)
No Information Rate : 0.71
P-Value [Acc > NIR] : 0.2954

Kappa : 0.3091
McNemar's Test P-Value : 0.2207

Sensitivity : 0.4483
Specificity : 0.8451
Pos Pred Value : 0.5417
Neg Pred Value : 0.7895
Prevalence : 0.2900
Detection Rate : 0.1300
Detection Prevalence : 0.2400
Balanced Accuracy : 0.6467

'Positive' Class : BAD

```
Acc=round(cM$overall[1],digits=2)
Sens=round(cM$byClass[1],digits=2)
Spec=cM$byClass[2]
```

The coefficients are not all important as evident from the diagnostics, which means the model could be improved via subset selection or some other method. In the above example the **Threshold** for selecting one of the states (i.e. classifying) of the response **0.5**. The accuracy of the model was **0.73** which is pretty good. If we did not know any more specifics about the problem at hand this would be effective. In the case that the bad valued differently the various errors, for example a “bad” creditor given a loan is more costly than a “good” one being denied a loan, we would like to check the sensitivity and specificity of the model. In this case the model has sensitivity of **0.45**, which is dependent on the True positive and False negative errors (“bad” creditors classified as “good”). Based on this the model does not perform well and we shall see in the following part how to optimize it.


Part 2: Given that the **false negative errors** (“bad” creditors classified as “good”) is 5 times more costly

than the false positive errors (“good” creditors identified as “bad”) we would like to have a our model to minimize the **Total cost function**. In this case the cost function is:

$$Cost = FN * 5 * C + FP * C$$

where C= positive constant. We could have plotted a **ROC** curve and select an appropriate cut-off point, but since we have the specific cost function I have decided to follow a different approach. This is equivalent to saying that we would like a model with higher **sensitivity**. In the following code I tried to create a simple optimization with a “for” loop, rather than use a premade package. For this task, I trained the model with 80% of the data, since the model remains the same for different Thresholds, calculated the cost function based on the performance on the test set for various Thresholds and selected the model which **minimizes** the total cost. I do not think there is need for validation set here, since the model remains the same.

```
Cost=matrix(0L,nrow=10,ncol=1)
range=seq(0.05,0.5,0.05)
for (i in 1:10){
  Pred2=ifelse(predicted>range[i], 'BAD', "GOOD")
  temp=confusionMatrix(table(Pred2,Actual[g=="test"]))
  Cost[i]=temp$table[2,1]*5+temp$table[1,2]
}
plot(range,Cost,xlab="Threshold value",ylab="Cost",main="Cost vs Threshold",type="o")
```



```
OptTh=range[which(Cost==min(Cost))]
final=confusionMatrix(table(ifelse(predicted>OptTh,
                                   'BAD', "GOOD"),Actual[g=="test"]))
Sens=round(final$byClass[1],digits=2)
```


In the above example I used various lower Thresholds in order to increase the sensitivity of the model. The **minimum Cost** was found for **Threshold value equal to 0.1**. In this case the sensitivity of the model was **0.93**, which is pretty good.